

Chest X-Ray Images (Pneumonia)

- Mayank Srivastava



- (<https://www.linkedin.com/in/mayank-srivastava-6a8421105/>)



Dataset: <https://www.kaggle.com/datasets/paulthothymooney/chest-xray-pneumonia> (<https://www.kaggle.com/datasets/paulthothymooney/chest-xray-pneumonia>)

Content

- The dataset is organized into 3 folders (train, test, val) and contains subfolders for each image category (Pneumonia/Normal).
- There are 5,863 X-Ray images (JPEG) and 2 categories (Pneumonia/Normal).

Source

- Chest X-ray images (anterior-posterior) were selected from retrospective cohorts of pediatric patients of one to five years old from Guangzhou Women and Children’s Medical Center, Guangzhou.
- All chest X-ray imaging was performed as part of patients’ routine clinical care.
- For the analysis of chest x-ray images, all chest radiographs were initially screened for quality control by removing all low quality or unreadable scans.
- The diagnoses for the images were then graded by two expert physicians before being cleared for training the AI system. - - In order to account for any grading errors, the evaluation set was also checked by a third expert.

Project Overview

- This is a Classification Problem.
- **Objective:** To classify chest X-ray images into two categories: Pneumonia and Normal using a CNN model.

Importing Libraries

```
In [1]: 1 import tensorflow as tf
2 from tensorflow import keras
3 from tensorflow.keras.models import Sequential
4 from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense, Dropout
5 from tensorflow.keras.preprocessing.image import ImageDataGenerator
6 from tensorflow.keras.optimizers import Adam # You can try different optimizers

In [2]: 1 import pandas as pd
2 import numpy as np
3 import matplotlib.pyplot as plt
4 import seaborn as sns
5
6 pd.set_option('display.max_rows', None)
7 pd.set_option('display.max_columns', None)
8
9 from sklearn.metrics import recall_score, confusion_matrix, roc_auc_score, roc_curve, auc
```

Setting the direcorey for train, test, val image data sets

```
In [3]: 1 train_dir = r'E:\DS journey\Deep Learning Datasets\chest_xray\train'
2 test_dir = r'E:\DS journey\Deep Learning Datasets\chest_xray\test'
3 val_dir = r'E:\DS journey\Deep Learning Datasets\chest_xray\val'

In [97]: 1 # Data Distribution
2
3 data=pd.DataFrame({
4     "Train":{i:len(os.listdir(os.path.join(train_dir,i))) for i in os.listdir(train_dir)},
5     "Test":{i:len(os.listdir(os.path.join(test_dir,i))) for i in os.listdir(test_dir)},
6 })
7 data
```

Out[97]:

| | Train | Test |
|-----------|-------|------|
| NORMAL | 1341 | 234 |
| PNEUMONIA | 3875 | 390 |

```
In [95]: 1 list(data.index)

Out[95]: ['NORMAL', 'PNEUMONIA']

In [98]: 1 [sum(data[i]) for i in data ]

Out[98]: [5216, 624]

In [114]: 1 sum(data.loc['NORMAL',:])

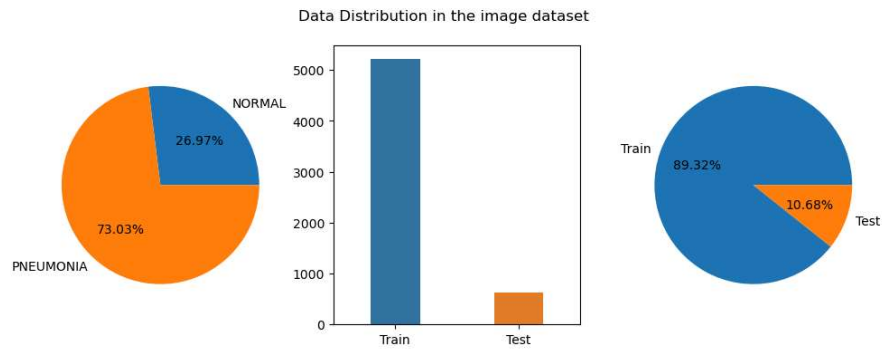
Out[114]: 1575

In [115]: 1 sum(data.loc['PNEUMONIA',:])

Out[115]: 4265
```

Data Distribution

```
In [123]: 1 plt.figure(figsize=(12,4))
2 plt.subplot(1,3,1)
3 plt.pie([sum(data.loc[i,:]) for i in list(data.index)], labels=data.index, autopct="%1.2f%%")
4
5
6 plt.subplot(1,3,2)
7 sns.barplot(x=data.columns, y=sum(data.values), width=0.4)
8
9 plt.subplot(1,3,3)
10 plt.pie([sum(data[i]) for i in list(data.columns)], labels=data.columns, autopct="%1.2f%%")
11
12 plt.suptitle('Data Distribution in the image dataset')
13 plt.show()
```



Setting seed for Reproducibility

```
In [4]: 1 # Set random seeds for reproducibility
2 import random
3 seed = 42
4 tf.random.set_seed(seed) # Sets the random seed for TensorFlow operations.
5 np.random.seed(seed) # Sets the random seed for NumPy operations.
6 random.seed(seed) # Sets the random seed for Python's built-in random module.
7
```

```
In [26]: 1 img_width, img_height = 240, 240 # Target image size
2 batch_size = 32
```

```
In [6]: 1 # Data generators with augmentation
2
3
4 # rotation_range=40,
5 # width_shift_range=0.2,
6 # height_shift_range=0.2,
7 # shear_range=0.2,
8 # zoom_range=0.2,
9 # horizontal_flip=True
```

Creating the ImageDataGenerator instance for test and train

```
In [7]: 1 train_datagen = ImageDataGenerator(rescale=1./255,
2                                     rotation_range=40,
3                                     width_shift_range=0.2,
4                                     height_shift_range=0.2,
5                                     shear_range=0.2,
6                                     zoom_range=0.2,
7                                     horizontal_flip=True)
8 test_datagen = ImageDataGenerator(rescale=1./255)
9 val_datagen = ImageDataGenerator(rescale=1./255)
```

```
In [8]: 1 train_generator = train_datagen.flow_from_directory(
2     train_dir,
3     target_size=(240, 240),
4     batch_size=32,
5     class_mode='binary') # Binary classification: cat or dog
```

Found 5216 images belonging to 2 classes.

Loading the image data from local direcotry

```
In [9]: 1 test_generator = test_datagen.flow_from_directory(
2     test_dir,
3     target_size=(240,240),
4     batch_size=32,
5     class_mode='binary') # Binary classification: cat or dog
```

Found 624 images belonging to 2 classes.

```
In [10]: 1 val_generator = val_datagen.flow_from_directory(
2     val_dir,
3     target_size=(240,240),
4     batch_size=batch_size,
5     class_mode='binary') # Binary classification: cat or dog
```

Found 16 images belonging to 2 classes.

CNN Model Architecture

```
In [11]: 1 model = Sequential()
2
3 model.add(Conv2D(32, (3, 3), activation='relu', input_shape=(240, 240, 3)))
4 model.add(MaxPooling2D((2, 2)))
5
6 model.add(Conv2D(64, (3, 3), activation='relu'))
7 model.add(MaxPooling2D((2, 2)))
8
9 model.add(Conv2D(128, (3, 3), activation='relu'))
10 model.add(MaxPooling2D((2, 2)))
11
12 model.add(Flatten())
13
14 ## ANN
15 model.add(Dense(128, activation='relu'))
16
17 model.add(Dense(64, activation='relu')) # hidden
18
19 model.add(Dropout(0.5)) # to prevent overfitting
20 model.add(Dense(1, activation='sigmoid')) # Output Layer ('NORMAL' or 'PNEUMONIA')
21
E:\anaconda3\Lib\site-packages\keras\src\layers\convolutional\base_conv.py:107: UserWarning: Do not pass an 'input_shape'/'input_dim' argument to a layer. When using Sequential models,
prefer using an 'Input(shape)' object as the first layer in the model instead.
  super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

```
In [12]: 1 model.compile(loss='binary_crossentropy',
2 optimizer='adam',
3 metrics=['accuracy'])
```

```
In [71]: 1 model.summary()
```

Model: "sequential"

| Layer (type) | Output Shape | Param # |
|--------------------------------|----------------------|------------|
| conv2d (Conv2D) | (None, 238, 238, 32) | 896 |
| max_pooling2d (MaxPooling2D) | (None, 119, 119, 32) | 0 |
| conv2d_1 (Conv2D) | (None, 117, 117, 64) | 18,496 |
| max_pooling2d_1 (MaxPooling2D) | (None, 58, 58, 64) | 0 |
| conv2d_2 (Conv2D) | (None, 56, 56, 128) | 73,856 |
| max_pooling2d_2 (MaxPooling2D) | (None, 28, 28, 128) | 0 |
| flatten (Flatten) | (None, 100352) | 0 |
| dense (Dense) | (None, 128) | 12,845,184 |
| dense_1 (Dense) | (None, 64) | 8,256 |
| dropout (Dropout) | (None, 64) | 0 |
| dense_2 (Dense) | (None, 1) | 65 |

Total params: 38,840,261 (148.16 MB)

Trainable params: 12,946,753 (49.39 MB)

Non-trainable params: 0 (0.00 B)

Optimizer params: 25,893,508 (98.78 MB)

```
In [13]: 1 from tensorflow.keras.callbacks import ModelCheckpoint
2 checkpoint = ModelCheckpoint(r'best_model.keras',
3                             monitor='val_accuracy', save_best_only=True, mode='max', verbose=1)
4
5
6 history = model.fit(train_generator, validation_data=test_generator, epochs = 10, callbacks=[checkpoint],
7                     shuffle = False)
8
9 # Load the best weights
10 model.load_weights('best_model.keras')
```

```
163/163 — 0s 4s/step - accuracy: 0.8935 - loss: 0.2578
Epoch 6: val_accuracy did not improve from 0.87500
163/163 — 742s 5s/step - accuracy: 0.8935 - loss: 0.2578 - val_accuracy: 0.8189 - val_loss: 0.3797
Epoch 7/10
163/163 — 0s 5s/step - accuracy: 0.8862 - loss: 0.2742
Epoch 7: val_accuracy did not improve from 0.87500
163/163 — 809s 5s/step - accuracy: 0.8863 - loss: 0.2740 - val_accuracy: 0.8734 - val_loss: 0.3192
Epoch 8/10
163/163 — 0s 5s/step - accuracy: 0.9121 - loss: 0.2163
Epoch 8: val_accuracy improved from 0.87500 to 0.88141, saving model to best_model.keras
163/163 — 786s 5s/step - accuracy: 0.9121 - loss: 0.2163 - val_accuracy: 0.8814 - val_loss: 0.3084
Epoch 9/10
163/163 — 0s 6s/step - accuracy: 0.9086 - loss: 0.2203
Epoch 9: val_accuracy did not improve from 0.88141
163/163 — 1028s 6s/step - accuracy: 0.9086 - loss: 0.2203 - val_accuracy: 0.8782 - val_loss: 0.3282
Epoch 10/10
163/163 — 0s 5s/step - accuracy: 0.9177 - loss: 0.2044
Epoch 10: val_accuracy did not improve from 0.88141
163/163 — 820s 5s/step - accuracy: 0.9177 - loss: 0.2044 - val_accuracy: 0.8365 - val_loss: 0.4421
```

Saving the model

```
In [14]: 1 model.save('my_model.keras')
```

```
In [15]: 1 new_model=keras.models.load_model('my_model.keras')
```

Model-evaluation

```
In [16]: 1 new_model.evaluate(test_generator)
```

20/20 — 34s 2s/step - accuracy: 0.8975 - loss: 0.2990

Out[16]: [0.308383047580719, 0.8814102411270142]

```
In [17]: 1 new_model.evaluate(train_generator)
```

163/163 — 487s 3s/step - accuracy: 0.9206 - loss: 0.1797

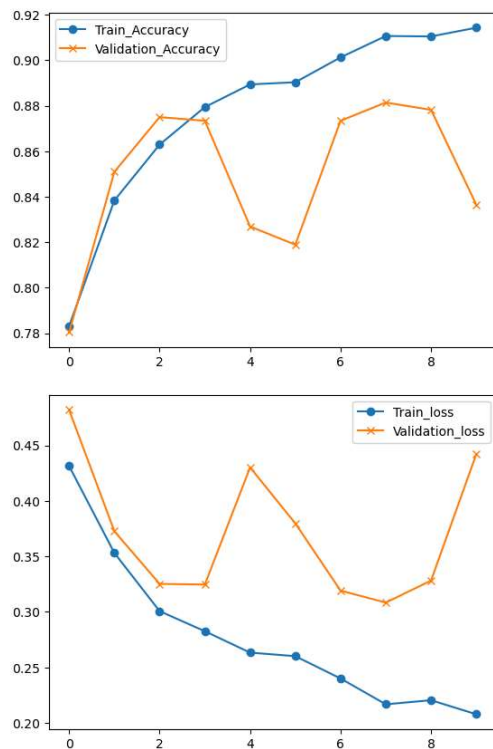
Out[17]: [0.18747134506702423, 0.9177530407905579]

```
In [18]: 1 new_model.evaluate(val_generator)
```

1/1 — 1s 1s/step - accuracy: 0.5625 - loss: 0.8518

Out[18]: [0.8517850041389465, 0.5625]

```
In [22]: # Plot accuracy and loss
1 import matplotlib.pyplot as plt
2 plt.plot(history.history['accuracy'], marker='o', label='Train_Accuracy')
3 plt.plot(history.history['val_accuracy'], marker='x', label='Validation_Accuracy')
4 plt.legend()
5 plt.show()
6
7 # Similar plot for loss
8 plt.plot(history.history['loss'], marker='o', label='Train_loss')
9 plt.plot(history.history['val_loss'], marker='x', label='Validation_loss')
10 plt.legend()
11 plt.show()
12
```



Make Predictions

```
In [53]: import os
1 # Load and preprocess a single image
2 img_paths = [os.path.join(r'E:\DS journey\Deep Learning Datasets\chest_xray\train\PNEUMONIA', i) \
3               for i in os.listdir(r'E:\DS journey\Deep Learning Datasets\chest_xray\train\PNEUMONIA')]
4 len(img_paths)
```

Out[53]: 3875

```
In [68]: 1 img_path = img_paths[2500:2550] # enter indices here
2 for img_path in img_path_:
3
4     # Load and preprocess a single image
5     img = tf.keras.preprocessing.image.load_img(img_path, target_size=(img_height, img_width))
6     img_array = tf.keras.preprocessing.image.img_to_array(img)
7     img_array = np.expand_dims(img_array, axis=0) / 255.0
8     #plt.imshow(img_array.reshape(240,240,3))
9     # Predict
10    prediction = model.predict(img_array)
11    print(prediction)
12    print('NORMAL' if prediction < 0.5 else 'PNEUMONIA')
13
```

```
[[0.0001201]]
PNEUMONIA
1/1 ----- 0s 337ms/step
[[0.9806409]]
PNEUMONIA
1/1 ----- 0s 457ms/step
[[0.9757314]]
PNEUMONIA
1/1 ----- 0s 189ms/step
[[0.8217449]]
PNEUMONIA
1/1 ----- 0s 373ms/step
[[0.9924891]]
PNEUMONIA
1/1 ----- 0s 447ms/step
[[0.74353504]]
PNEUMONIA
1/1 ----- 0s 293ms/step
[[0.9882326]]
PNEUMONIA
```

Observations:

1. This is an Image Classification Problem
2. The CNN Model takes the Xray scan image as input and predicts if the patient is NORMAL or infected with PNEUMONIA
3. The Model performance is as follows:
 - training prediction accuracy of 92.06% with loss of 0.1797 and
 - test prediction accuracy of 89.75% with loss of 0.2990
4. Accuracy can be further improved by tuning the model parameters as:
 - optimizer
 - learning_rate
 - no. of convolution and pooling layers
 - no. of Dense layers

