

MNIST digits classification using CNN

-Mayank Srivastava



About Dataset

Source

The MNIST database (Modified National Institute of Standards and Technology database) is a large database of handwritten digits that is commonly used for training various image processing systems.

Dataset link: <https://www.kaggle.com/datasets/scolianni/mnistasjpg> (<https://www.kaggle.com/datasets/scolianni/mnistasjpg>)

Objective:

To classify images of handwritten digits (0-9) using a CNN model.

Summary

- Data Preparation: Created TRAIN and TEST directories and its sub-folders (classes). Using shutil, moved files in the ratio 80-20 respectively for each class.
- Model Building: Defined a CNN architecture suitable for digit recognition.
- Model Training: Trained the model using the training data.
- Evaluation: Evaluated the model on the test data and plotted the training history.
- Prediction: Made predictions on new images using the trained model.

Importing Libraries

```
In [1]: 1 import pandas as pd
2 import numpy as np
3 import matplotlib.pyplot as plt
4 import seaborn as sns
5
6 # pd.set_option('display.max_rows', None)
7 # pd.set_option('display.max_columns', None)
8

In [2]: 1 import tensorflow as tf
2 from tensorflow import keras
3 from tensorflow.keras.models import Sequential
4 from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense, Dropout
5 from tensorflow.keras.preprocessing.image import ImageDataGenerator
6 from tensorflow.keras.optimizers import Adam # You can try different optimizers
7
8
```

Setting the directory for train & test image data sets

```
In [3]: 1 source= r'E:\DS journey\Deep Learning Datasets\Digit Image Dataset\Main set'
2 train_dir = r'E:\DS journey\Deep Learning Datasets\Digit Image Dataset\TRAIN'
3 test_dir = r'E:\DS journey\Deep Learning Datasets\Digit Image Dataset\TEST'

In [4]: 1 import shutil
2 import os
3 from sklearn.model_selection import train_test_split

In [5]: 1 classes =[i for i in os.listdir(source)]
2 classes

Out[5]: ['0', '1', '2', '3', '4', '5', '6', '7', '8', '9']
```

Creating sub-folders for all classes in TRAIN and TEST

```
In [6]: 1 ## creating directories in TRAIN & TEST, based in Main set
2 ## This code needs to be executed only once
3
4
5 # for i in classes:
6 #     os.mkdir(os.path.join(train_dir,i))
7 #     os.mkdir(os.path.join(test_dir,i))
```

Moving files based on 80-20 split

python code for traversing all classes of Source file

```
for i in os.listdir(source):
    j= os.path.join(source,i)
    print(j)

C:\Users\hp\Python Datasets\Deep Learning\CNN\Digit_Recognizer_CNN\archive (3)\Main set\0
C:\Users\hp\Python Datasets\Deep Learning\CNN\Digit_Recognizer_CNN\archive (3)\Main set\1
C:\Users\hp\Python Datasets\Deep Learning\CNN\Digit_Recognizer_CNN\archive (3)\Main set\2
C:\Users\hp\Python Datasets\Deep Learning\CNN\Digit_Recognizer_CNN\archive (3)\Main set\3
C:\Users\hp\Python Datasets\Deep Learning\CNN\Digit_Recognizer_CNN\archive (3)\Main set\4
C:\Users\hp\Python Datasets\Deep Learning\CNN\Digit_Recognizer_CNN\archive (3)\Main set\5
C:\Users\hp\Python Datasets\Deep Learning\CNN\Digit_Recognizer_CNN\archive (3)\Main set\6
C:\Users\hp\Python Datasets\Deep Learning\CNN\Digit_Recognizer_CNN\archive (3)\Main set\7
C:\Users\hp\Python Datasets\Deep Learning\CNN\Digit_Recognizer_CNN\archive (3)\Main set\8
C:\Users\hp\Python Datasets\Deep Learning\CNN\Digit_Recognizer_CNN\archive (3)\Main set\9
```

```
In [7]: 1 # for i in os.listdir(source):
2 #     j= os.path.join(source,i)
3
4 #     all_images = os.listdir(j)
5 #     train_images, test_images =train_test_split(all_images, test_size =0.2, random_state =42)
6
7 #     for img in train_images:
8 #         shutil.move(os.path.join(source,i, img), os.path.join(train_dir,i,img))
9
10 #     for img in test_images:
11 #         shutil.move(os.path.join(source,i, img), os.path.join(test_dir,i,img))
```

Creating ImageDataGenerator objects

```
In [8]: 1 train_datagen=ImageDataGenerator(rescale=1./255,
2                                     rotation_range=40,
3                                     width_shift_range =0.2,
4                                     height_shift_range =0.2,
5                                     shear_range =0.2,
6                                     zoom_range =0.2,
7                                     horizontal_flip = True,)
8 test_datagen= ImageDataGenerator(rescale =1/255)
```

Loading TRAIN & TEST images using ImageDataGenerator objects

```
In [9]: 1 train_generator = train_datagen.flow_from_directory(train_dir,
2                                                         target_size= (128,128),
3                                                         batch_size =32,
4                                                         class_mode ='categorical')
5 test_generator =test_datagen.flow_from_directory(test_dir,
6                                                  target_size =(128,128),
7                                                  batch_size =32,
8                                                  class_mode ='categorical')
```

Found 33595 images belonging to 10 classes.
Found 8405 images belonging to 10 classes.

```
In [10]: 1 d = {"Classes":[i for i in os.listdir(source)],
2         "Train_images":[len(os.listdir(os.path.join(train_dir,i))) for i in os.listdir(train_dir)],
3         "Test_images":[len(os.listdir(os.path.join(test_dir,i))) for i in os.listdir(test_dir)]
4     }
5
6 df= pd.DataFrame(d)
7 df
```

```
Out[10]:
```

	Classes	Train_images	Test_images
0	0	3305	827
1	1	3747	937
2	2	3341	836
3	3	3480	871
4	4	3257	815
5	5	3036	759
6	6	3309	828
7	7	3520	881
8	8	3250	813
9	9	3350	838

```
In [11]: 1 # Total No. of images
2 sum(df.Train_images)+sum(df.Test_images)
```

```
Out[11]: 42000
```

```
In [12]: 1 print("Total Training images: ",sum(df.Train_images),"Total Testing images: ", sum(df.Test_images))
```

Total Training images: 33595
Total Testing images: 8405

Creating ImageDataGenerator objects

```
In [13]: 1 train_generator = train_datagen.flow_from_directory(
2         train_dir,
3         target_size=(128, 128),
4         batch_size=32,
5         class_mode='categorical') # Binary classification: cat or dog
```

Found 33595 images belonging to 10 classes.

```
In [14]: 1 test_generator = test_datagen.flow_from_directory(
2         test_dir,
3         target_size=(128,128),
4         batch_size=32,
5         class_mode='categorical') # Binary classification: cat or dog
```

Found 8405 images belonging to 10 classes.

```
In [15]: 1 num_classes=len(train_generator.class_indices)
2 train_generator.class_indices
```

```
Out[15]: {'0': 0,
'1': 1,
'2': 2,
'3': 3,
'4': 4,
'5': 5,
'6': 6,
'7': 7,
'8': 8,
'9': 9}
```

CNN_Model_Architecture

```
In [16]: M 1 # CNN
2 model = Sequential()
3
4 model.add(Conv2D(32, (3, 3), activation='relu', input_shape=(128, 128, 3)))
5 model.add(MaxPooling2D((2, 2)))
6
7 model.add(Conv2D(64, (3, 3), activation='relu', input_shape=(128, 128, 3)))
8 model.add(MaxPooling2D((2, 2)))
9 model.add(Conv2D(128, (3, 3), activation='relu', input_shape=(128, 128, 3)))
10 model.add(MaxPooling2D((2, 2)))
11 model.add(Flatten())
12
13 ## ANN
14 model.add(Dense(100, activation='relu'))
15
16 model.add(Dense(50, activation='relu')) # hidden
17
18 model.add(Dropout(0.5)) # Regularization to prevent overfitting
19 model.add(Dense(num_classes, activation='softmax')) # Output layer (cat or dog)
```

E:\anaconda3\Lib\site-packages\keras\src\layers\convolutional\base_conv.py:107: UserWarning: Do not pass an 'input_shape'/'input_dim' argument to a layer. When using Sequential models, prefer using a 'Input(shape)' object as the first layer in the model instead.
super().__init__(activity_regularizer=activity_regularizer, **kwargs)

```
In [17]: M 1 # Compiling the model
2 model.compile(
3     optimizer='adam',
4     loss='categorical_crossentropy', # Use categorical_crossentropy for multi-class classification
5     metrics=['accuracy']
6 )
7
8 # Printing model summary
9 model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 126, 126, 32)	896
max_pooling2d (MaxPooling2D)	(None, 63, 63, 32)	0
conv2d_1 (Conv2D)	(None, 61, 61, 64)	18,496
max_pooling2d_1 (MaxPooling2D)	(None, 30, 30, 64)	0
conv2d_2 (Conv2D)	(None, 28, 28, 128)	73,856
max_pooling2d_2 (MaxPooling2D)	(None, 14, 14, 128)	0
flatten (Flatten)	(None, 25088)	0
dense (Dense)	(None, 100)	2,508,900
dense_1 (Dense)	(None, 50)	5,050
dropout (Dropout)	(None, 50)	0
dense_2 (Dense)	(None, 10)	510

Total params: 2,607,708 (9.95 MB)

Trainable params: 2,607,708 (9.95 MB)

Non-trainable params: 0 (0.00 B)

Using ModelCheckpoint Callback

```
In [18]: M 1 # Define ModelCheckpoint callback to save the best weights
2 from tensorflow.keras.callbacks import ModelCheckpoint
3 checkpoint = ModelCheckpoint(r'best_model.keras',
4                             monitor='val_accuracy', save_best_only=True, mode='max', verbose=1)
5
6 # Train the model
7 history = model.fit(
8     train_generator,
9     epochs=10,
10    validation_data=test_generator,
11    callbacks=[checkpoint],
12    shuffle = False
13 )
14
15
```

E:\anaconda3\Lib\site-packages\keras\src\trainers\data_adapters\py_dataset_adapter.py:121: UserWarning: Your 'PyDataset' class should call 'super().__init__(**kwargs)' in its construct or: '**kwargs' can include 'workers', 'use_multiprocessing', 'max_queue_size'. Do not pass these arguments to 'fit()', as they will be ignored.
self._warn_if_super_not_called()

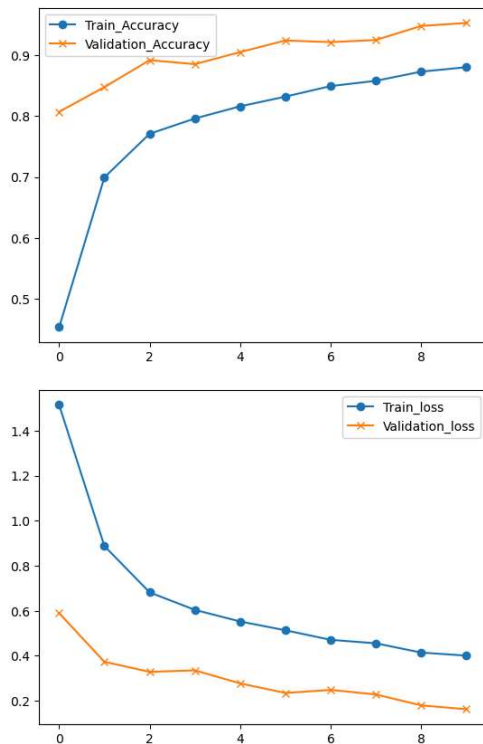
1050/1050 — 0s 695ms/step - accuracy: 0.3104 - loss: 1.8700
Epoch 1: val_accuracy improved from -inf to 0.80690, saving model to best_model.keras
1050/1050 — 813s 772ms/step - accuracy: 0.3106 - loss: 1.8697 - val_accuracy: 0.8069 - val_loss: 0.5895
Epoch 2/10
1050/1050 — 0s 484ms/step - accuracy: 0.6693 - loss: 0.9725
Epoch 2: val_accuracy improved from 0.80690 to 0.84735, saving model to best_model.keras
1050/1050 — 584s 556ms/step - accuracy: 0.6694 - loss: 0.9724 - val_accuracy: 0.8474 - val_loss: 0.3735
Epoch 3/10
1050/1050 — 0s 618ms/step - accuracy: 0.7658 - loss: 0.7009
Epoch 3: val_accuracy improved from 0.84735 to 0.89149, saving model to best_model.keras
1050/1050 — 676s 644ms/step - accuracy: 0.7658 - loss: 0.7009 - val_accuracy: 0.8915 - val_loss: 0.3288
Epoch 4/10
1050/1050 — 0s 426ms/step - accuracy: 0.7883 - loss: 0.6280
Epoch 4: val_accuracy did not improve from 0.89149
1050/1050 — 474s 452ms/step - accuracy: 0.7883 - loss: 0.6279 - val_accuracy: 0.8849 - val_loss: 0.3355
Epoch 5/10
1050/1050 — 0s 429ms/step - accuracy: 0.8118 - loss: 0.5656
Epoch 5: val_accuracy improved from 0.89149 to 0.90470, saving model to best_model.keras
1050/1050 — 478s 455ms/step - accuracy: 0.8118 - loss: 0.5656 - val_accuracy: 0.9047 - val_loss: 0.2776
Epoch 6/10
1050/1050 — 0s 426ms/step - accuracy: 0.8311 - loss: 0.5115
Epoch 6: val_accuracy improved from 0.90470 to 0.92374, saving model to best_model.keras
1050/1050 — 476s 453ms/step - accuracy: 0.8311 - loss: 0.5115 - val_accuracy: 0.9237 - val_loss: 0.2350
Epoch 7/10
1050/1050 — 0s 427ms/step - accuracy: 0.8488 - loss: 0.4705
Epoch 7: val_accuracy did not improve from 0.92374
1050/1050 — 475s 453ms/step - accuracy: 0.8488 - loss: 0.4705 - val_accuracy: 0.9210 - val_loss: 0.2487
Epoch 8/10
1050/1050 — 0s 438ms/step - accuracy: 0.8521 - loss: 0.4652
Epoch 8: val_accuracy improved from 0.92374 to 0.92457, saving model to best_model.keras
1050/1050 — 488s 465ms/step - accuracy: 0.8521 - loss: 0.4652 - val_accuracy: 0.9246 - val_loss: 0.2285
Epoch 9/10
1050/1050 — 0s 428ms/step - accuracy: 0.8689 - loss: 0.4252
Epoch 9: val_accuracy improved from 0.92457 to 0.94753, saving model to best_model.keras
1050/1050 — 477s 454ms/step - accuracy: 0.8689 - loss: 0.4252 - val_accuracy: 0.9475 - val_loss: 0.1799
Epoch 10/10
1050/1050 — 0s 428ms/step - accuracy: 0.8767 - loss: 0.4076
Epoch 10: val_accuracy improved from 0.94753 to 0.95241, saving model to best_model.keras
1050/1050 — 502s 454ms/step - accuracy: 0.8767 - loss: 0.4076 - val_accuracy: 0.9524 - val_loss: 0.1627

Model-evaluation

```
In [19]: 1 # Load the best weights
2 model.load_weights('best_model.keras')
3
4 # Evaluate the model
5 test_loss, test_accuracy = model.evaluate(test_generator)
6 print(f'Test accuracy: {test_accuracy:.4f}')
7 print(f'Test loss: {test_loss:.4f}')
8
9 # Evaluate the model
10 train_loss, train_accuracy = model.evaluate(train_generator)
11 print(f'Train accuracy: {train_accuracy:.4f}')
12 print(f'Train loss: {train_loss:.4f}')

263/263 ————— 28s 106ms/step - accuracy: 0.9516 - loss: 0.1566
Test accuracy: 0.9524
Test loss: 0.1627
1050/1050 ————— 261s 249ms/step - accuracy: 0.9062 - loss: 0.2954
Train accuracy: 0.9058
Train loss: 0.2925
```

```
In [20]: 1 import matplotlib.pyplot as plt
2 plt.plot(history.history['accuracy'], marker='o', label='Train_Accuracy')
3 plt.plot(history.history['val_accuracy'], marker='x', label='Validation_Accuracy')
4 plt.legend()
5 plt.show()
6
7 # Similar plot for loss
8 plt.plot(history.history['loss'], marker='o', label='Train_loss')
9 plt.plot(history.history['val_loss'], marker='x', label='Validation_loss')
10 plt.legend()
11 plt.show()
```



Make Predictions

```
In [21]: 1 import os
2 # Load and preprocess a single image
3 img_paths= [os.path.join(r'E:\DS journey\Deep Learning Datasets\Digit Image Dataset\testSet\testSet', i) for i in \
4              os.listdir(r'E:\DS journey\Deep Learning Datasets\Digit Image Dataset\testSet\testSet')]
5
6 len(img_paths)
```

Out[21]: 28000

```

In [22]: 1 img_path=img_paths[27010:27030] # enter indices here
2
3 for img_path in img_path_:
4
5     # Load and preprocess a single image
6     img = tf.keras.preprocessing.image.load_img(img_path, target_size=(128, 128))
7     img_array = tf.keras.preprocessing.image.img_to_array(img)
8     img_array = np.expand_dims(img_array, axis=0) / 255.0
9
10    # Predict
11    prediction = model.predict(img_array)
12    predicted_labels = int(np.argmax(prediction, axis=1).reshape(1,1))
13    # image plot part
14    plt.figure(figsize =(0.5,0.5))
15    plt.tick_params(left = False, right = False , labelleft = False ,
16                  labelbottom = False, bottom = False)
17
18    plt.imshow(img_array.reshape(128,128,3))
19    plt.show()
20    print("Value of the above image is:",predicted_labels)
21

```

1/1 ————— 0s 163ms/step



value of the above image is: 2

1/1 ————— 0s 33ms/step



value of the above image is: 3

1/1 ————— 0s 33ms/step



Observations

- The Digit Dataset has 10 classes from 0,1,2...9.
- Firstly created the TRAIN & TEST folders and sub-folders based on classes.
- Using the shutil library copied images from source to respective directories maintaining a 80-20 split.
- Loaded the image data using ImageDataGenerator class object and methods.
- CNN model was created and Results after 10 epochs with 3 Conv layers, 3 max pooling and 3 Dense layers is as follows:
 - Train accuracy: 0.9058- loss: 0.2925
 - Test accuracy: 0.9524- loss: 0.1627